

From **Black Box** to Inspectable AI

Future Thinking Analysis on Mechanistic Interpretability of LLM's

Elad Moshe

Emerging Technologies Course

MCSBT Program - IE University Madrid

January 2026



Future Thinking Analysis on Mechanistic Interpretability of LLM's

Background and Technology Overview

Mechanistic Interpretability (MI for short) seeks to understand the reasoning mechanisms of large language models. After the "gold rush" of improving performance, **the focus now shifts** to revealing how specific behaviors of the LLM's arise. By uncovering these once perceived "**Black Boxes**", researchers are trying to develop tools that will allow us to understand how LLM's internally represent and transform information.

LLMs' reasoning patterns and inner workings **still remain largely hidden**, with current understanding only scratching the surface. Truly explainable performant AI is still far from reach, but several advancements in the field of MI sparked a lot of hope among the AI community.

Growing global consumer reliance on AI makes the **explainability** and **alignment** of LLMs **essential** for safety, regulation and bias reduction. Beyond risk mitigation, interpretability enables **deeper understanding** and more effective use of AI in science and other domains—rather than simply beating chess grandmasters, we aim to understand how such performance is achieved.

Some examples of the tools that were developed in recent years by LLM interpretability researchers are techniques such as:

1. **Activation Patching** - tracing which activations cause specific outputs.
2. **Zeroing Components** - testing causal importance by changing/removing neurons
3. **Probing Classifiers** - assessing what information is encoded in internal model states.
4. **Attention and Logit Lenses** – inspecting outputs mid-run.
5. **Sparse Auto-Encoders** - isolating interpretable features from activations.
6. **Features Visualization** - Geometry analysis of embedding vectors representing different features.

These allow us to identify "Neural Circuits", "Chain of Thoughts", features location and **cause-effect pathways** inside models. Such newly discovered knowledge allows us to **enhance, suppress** or even, to some extent, **selectively turn on and off** behaviors and capabilities within the LLMs (without retraining).

Beside the tools that were being developed, the true breakthrough lies in the **empirical successes** showing that **meaningful internal mechanisms** can be isolated **even in large transformer models** – with "**The Golden Gate Bridge**" article by Anthropic¹ as the pivotal turning point.

This is considered a breakthrough now because LLMs just started to be **adopted daily by billions** around the world, the **models are bigger than ever** causing safety and alignment to be of the outmost importance. The LLM's that are being deployed still have very limited transparency. MI Tech addresses the broader problem of AI trust, safety and alignment by enabling governance scrutiny on the models' inner workings.

¹ See Bibliography *Templeton, 2024*

Futures Thinking Analysis

This section will showcase a future-thinking analysis of Interpretability Tech, it will focus on the current Signals, prominent Drivers and systematic implications this tech might have.

Signal	What it Indicates
Rising concern about alignment failures and hallucinations	Growing pressure to move beyond patch-based fixes toward deeper cause-effect behavioral understanding
Large investments by major AI players	Strategic shift from AI performance gains to explainability and safety
Interpretability is linked to safety and regulation	Future legitimacy of AI systems may depend on whether they can be inspected, audited and explained
Empirical breakthroughs revealing internal model structures	AI systems are now seen as more analyzable rather than pure black boxes (this may also create unrealistic expectations)
ML tools are Open-source and research is done public	Interpretability is moving from niche academic research toward broader adoption and standardized practice
Monetization Models in the LLM market are still unclear	Interpretability enables greater control over output, supporting value capture through personalization, data leverage or ability to influence End-Users
Fear of the AI market hype correction	Interpretability can help maintain expectations high until monetization models are clearer

Main drivers

The surge in **global AI hype** and **widespread adoption** amplifies both opportunity and risk. **LLMs are rapidly increasing in size and complexity**, challenging scalable alignment. Their commercial deployment in **high-stakes domains** such as healthcare, finance and defense intensifies **regulatory** and **public pressure** for transparency and accountability, especially around **user data** and **truthfulness**. Current surface-level alignment methods are reaching their limits, while **competition between AI labs** and **beyond-human performance** further drives the need to understand how and why models work internally.

1st and 2nd order effects

In the short term, mechanistic interpretability can lead to more **robust alignment**, reduced "**jail-breakability**" and increased **trust** among users and regulators, at the cost of **slower** and **more expensive deployment**.

Over time, this may **redistribute power** to organizations focusing on transparency rather than performance, **raise public expectations** for **accountability** and encourage regulators to rely on technical interpretability standards. This might also prove **useful in medical** and **scientific** discoveries due to better explainability and accountability.

Social Impact Analysis- Advanced Intelligence for all, Truth is Optional...

As with Social Networks, News Agencies and digital reliance → MI has the potential to democratize expert-level knowledge world-wide. While at the same time having catastrophic social impacts. Engagement-optimized deployment can distort truth, weaken trust and damage users' overall well-being.

The immediate beneficiaries of this tech are **regulators, policymakers** and the **tech providers** themselves. By auditing models behavior, they can support **evidence-based governance** and tackle misalignment better. This tech potentially allows for a more **factual-oriented outputs**, reducing hallucinations and misinformation.

MI is not only beneficial to those who seek a more secure and aligned AI system, but also to those who learn and gather insights from the models' **internal reasoning mechanisms** → many of them could be completely unknown and beneficial².

Scientists and **medical researchers** could inspect and critique the reasoning behind AI-generated decisions or hypotheses, allowing outputs to function as contestable scientific inputs rather than just help productivity³.

Public institutions and **high-stakes** sectors such as **healthcare, finance, law** and **education** may further benefit from more inspectable LLMs, indirectly improving outcomes for end users by being able to follow the **actual reasoning** behind outputs and have actual accountability. As a matter of fact, LLMs lack of transparency is currently **one of the main obstacles to wide-spread adoption in these fields**.

Smaller firms, startups, and **open-source communities** have struggled to compete in the costly AI performance race against Big Tech, but this recent technology may **lower barriers to entry** by shifting advantage toward those who master it.

Those new capabilities raise questions about who should have access to model internals, who is authorized to interpret them, and how—or whether—they should be regulated. As with **social media**, the widespread use of commercial LLMs **risks being shaped by incentives to maximize engagement and profit** over safety. Potentially manipulating users to stay on their screens, increasing **polarization** and **social division** by enabling the spread of **misinformation** or **harmful narratives**.

Imagine an LLM that **rewards user satisfaction over accuracy**, such model could present confident but incorrect answers or selectively frame facts to match user's expectations, gradually eroding trust in the objective truth, creating a new form of narrative manipulation. LLMs could also be tempered with in order to impose political narratives directly on users.

On the other hand, if this tech won't live up to technical expectations and won't be able to fully interpret the LLM's internal reasoning, partial explanations **may foster overconfidence**, creating a **false sense of safety or legitimacy**.

² In Douglas Adams "The Hitchhiker's Guide to the Galaxy", an advanced civilization builds a city-sized supercomputer to answer the "Ultimate Question of Life, the Universe, and Everything." After millions of years, it outputs the answer "42". This comic failure to understand the machine's internal reasoning mirrors the exact problem we face today with LLMs.

³ More research even point out that MI could one day lay the foundations to assess whether an AGI model is sentient or not - see Bibliography *AI Mashai 2025*

Reflective Analysis of Uncertainty

Main Areas of Uncertainty		
Will Tech Deliver? Will MI be able to fully explain LLMs behaviors, high-level concepts and reasoning?	Scalability How scalable MI will be within increasingly large, multi-model environments?	From Safety to Control Whether the technology can move beyond explainability toward meaningful behavioral control of the LLMs on a large scale
Potential Obstacles		
Lack of regulation & Oversight Lack of oversight reduces incentives for adoption, enables misuse and weakens trust in the tech's mission statement.	Reliance on Expert Knowledge Understanding MI insights could require reliance on experts across multiple specialized fields – hurting scalability	Tools Complexity High tool complexity may limit open-source adoption and scaling ⁴ .
Key Risk		Tensions
Malicious Use MI tools could be exploited for hacking, fraud and spreading misinformation by revealing exploitable model behaviors and allowing complex Jailbreaking.		Safety vs TTM and performance Tensions between safety, performance and TTM may discourage investment in slower more interpretable systems or affect overall performance.
Unintended Consequences		
Threat of Imitation MI may unintentionally enable imitation by revealing key mechanisms of LLMS, allowing competitors to replicate core capabilities without equivalent investment. For example, using disclosed reasoning pathways to reproduce performance-critical components in rival models. This may weaken incentives to open-source over IP protection.		

To conclude, even if the AI market will undergo hype correction, I believe this technology is still likely to remain relevant. Incentives for explainability are distributed across the entire market ecosystem, resources remain sufficient and empirical results continue to be promising.

Overall, the rise of MI signals a shift from performance-driven AI toward greater emphasis on transparency and safety. While it offers clear benefits for researchers and regulators, it also introduces a question about engagement vs truth-telling.

The tech has a potential to shift power toward actors with expertise and resources to interpret and act on MI insights. Significant uncertainty remains around ability to deliver, scalability and usability. Overall impact will eventually be determined by governance choices and institutional adoption among the large AI players and end users.

⁴ Today's MI tools need deep mathematical understanding. That may risk concentrating power toward the handful of people who can grasp the concepts behind MI.

Bibliography

Academic Papers

1. Al Mashai, A. (2025). *Beyond mechanistic interpretability: A philosophical framework for interpreting AI world models ("Digital Ontos") toward operational, ontological, and governance-aware AI interpretation*. SSRN. <https://doi.org/10.2139/ssrn.5966555>
2. Gekhman, Z., Ben David, E., Orgad, H., Ofek, E., Belinkov, Y., Szpektor, I., Herzig, J., & Reichart, R. (2025). *Inside-out: Hidden factual knowledge in LLMs*. arXiv. <https://doi.org/10.48550/arXiv.2503.15299>
3. Gurnee, W., Ameisen, E., Kauvar, I., Tarng, J., Pearce, A., Olah, C., & Batson, J. (2025). *When models manipulate manifolds: The geometry of a counting task* (Technical report). Anthropic. <https://transformer-circuits.pub/2025/linebreaks/index.html>
4. Orgad, H., Toker, M., Gekhman, Z., Reichart, R., Szpektor, I., Kotek, H., & Belinkov, Y. (2025). *LLMs know more than they show*. In *Proceedings of the Thirteenth International Conference on Learning Representations (ICLR 2025)*. <https://openreview.net/forum?id=KRnsX5Em3W>
5. Power, A., Burda, Y., Edwards, H., Babuschkin, I., & Misra, V. (2022). *Grokking: Generalization beyond overfitting on small algorithmic datasets*. arXiv. <https://doi.org/10.48550/arXiv.2201.02177>
6. Saphra, N., & Wiegrefe, S. (2024). *Mechanistic?* In *Proceedings of the 7th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP* (pp. 480–498). Association for Computational Linguistics. <https://aclanthology.org/2024.blackboxnlp-1.30.pdf>
7. Templeton, A., Conerly, T., Marcus, J., Lindsey, J., Bricken, T., Chen, B., Pearce, A., Citro, C., Ameisen, E., Jones, A., Cunningham, H., Turner, N. L., McDougall, C., MacDiarmid, M., Freeman, C. D., Sumers, T. R., Rees, E., Batson, J., Jermyn, A., Carter, S., Olah, C., & Henighan, T. (2024). *Scaling monosemanticity: Extracting interpretable features from Claude 3 Sonnet*. *Transformer Circuits Thread*. <https://transformer-circuits.pub/2024/scaling-monosemanticity/index.html>

Blogs

1. Amodei, D. (2025, April). *The urgency of interpretability*. Dario Amodei (blog). <https://www.darioamodei.com/post/the-urgency-of-interpretability>
2. Anthropic. (2024, May 23). *Golden Gate Claude*. Anthropic. <https://www.anthropic.com/news/golden-gate-claude>
3. Nanda, N. (2024, July 7). *An extremely opinionated annotated list of my favourite mechanistic interpretability papers v2*. Neel Nanda (blog). <https://www.alignmentforum.org/posts/NfFST5Mio7BCAQHPA/an-extremely-opinionated-annotated-list-of-my-favourite>

Online Lectures and Interviews

1. International Association for Safe & Ethical AI. (n.d.). *An introduction to mechanistic interpretability* [Video]. YouTube. <https://www.youtube.com/watch?v=0704iLc55Fs>
2. NLP Programming Labs. (n.d.). *Mechanistic interpretability for NLP: One-stop guide for everything you need to know* [Video]. YouTube. <https://www.youtube.com/watch?v=WSjLp2CpGsA>
3. Welch Labs. (n.d.). *The dark matter of AI [Mechanistic interpretability]* [Video]. YouTube. https://www.youtube.com/watch?v=UGO_Ehywuxc

Journalism

1. Douglas Heaven, W. (2026, January 12). *Meet the new biologists treating LLMs like aliens*. *MIT Technology Review*. <https://www.technologyreview.com/2026/01/12/1129782/ai-large-language-models-biology-alien-autopsy/>